

When does dispatch fidelity matter for battery sizing?

A regime-classification diagnostic

Julian Quick

Lawrence Berkeley National Laboratory

May 8, 2026

Abstract

Tesla Autobidder, Fluence Mosaic, and Wärtsilä GEMS run battery dispatch on machine-learning-driven stochastic optimization. NREL REopt, DTU **hydesign**, and most academic hybrid-power-plant sizing tools embed deterministic LP inner dispatch. Should sizing tools change architecture? We propose a practitioner-runnable regime-classification diagnostic — the b_{sat}^ϵ overlap test — that any operator can compute on one year of price data to answer the question for their own market. The diagnostic returns “invariance survives” on synthetic AR(1), all three years of Danish DK1 (Energinet, including the 2022 European energy crisis with max 871/MWh), and all three years of ERCOT North Hub (gridstatus, including February 2021 Storm Uri with sustained \$9000/MWh price-cap hours). Despite stochastic dispatch yielding 9–34% realized-NPV uplift at the optimum, the optimal battery energy capacity b_E^* is invariant across deterministic vs. multi-lag-ensemble dispatch in every (market, year, cost-form) configuration tested. We sketch a theoretical proposition motivating the diagnostic, then test its prediction empirically: it predicted that multi-day spectral content should break invariance; the data falsify the prediction. The empirical answer is broader than the theory anticipated. We release the benchmark (**battery_gym**) and the diagnostic for replication on any market.

1 Introduction

A battery operator answers two questions: *what to build* (capacity) and *how to operate it* (dispatch under uncertain prices and renewables). Industry decomposes them: capacity is a one-shot investment, dispatch is continuous re-optimization. The two communities of practice make different modelling choices. Tesla Autobidder has dispatched > 100 GWh of energy in global electricity markets; the platform’s public description ([Tesla, 2024](#)) cites “classical statistics, machine learning and numerical optimization” — explicitly stochastic optimization with ML price forecasts. Fluence Mosaic (16 GW deployed/awarded) is similarly a “stochastic optimization” bidding system ([Fluence, 2024](#)). Wärtsilä GEMS Pulse markets predictive analytics for revenue resilience ([Wärtsilä, 2025](#)). Together these platforms cover tens of gigawatts.

Meanwhile, academic hybrid-power-plant sizing tools (**hydesign** ([Murcia León et al., 2024](#)), REopt ([NREL, 2025](#)), PyPSA capacity-expansion variants) approximate inner dispatch as deterministic LP fed point-forecast prices and resources. The outer surrogate optimizer evaluates this LP 10^3 – 10^4 times per run; making it stochastic would cost 100–1000× more compute per evaluation.

Are the sizing tools structurally wrong? A practitioner cannot answer this on first principles. Sometimes deterministic-LP and stochastic dispatch produce identical capacity recommendations; sometimes not. The right question is empirical and per-market.

Contribution. We propose a practitioner-runnable diagnostic: the b_{sat}^ϵ overlap test (§3.2), which any operator can compute on one year of historical price data to test whether their market falls in the regime where deterministic-LP-inner-dispatch is empirically robust. The test is cheap, falsifiable, and survives independently of any theoretical assumption.

We motivate the diagnostic with a theoretical sketch (§3), then test the diagnostic — and the theoretical prediction it scaffolds — on:

- Synthetic diurnal AR(1) prices.
- Danish DK1 day-ahead prices for 2021/2022/2023, including the 2022 European energy crisis (max 871/MWh, 168 h spectral peak).
- ERCOT HB_NORTH day-ahead prices for 2021/2022/2023, including February 2021 Storm Uri (sustained \$9000/MWh price-cap hours, kurtosis 165).

In every (market, year) the diagnostic returns “invariance survives”: the optimal capacity b_E^* is identical across deterministic and stochastic dispatch policies, even though stochastic dispatch yields 9–34% realized-NPV uplift at the optimum. The theoretical sketch predicted that multi-day spectral content should break invariance. It does not. We surface the falsification explicitly (§4), discuss candidate mechanisms (§5), and note that the diagnostic survives the falsification: the proposition’s job was to motivate the test, not to determine its outcome.

Implication for practice. Across the markets and regimes tested, deterministic-LP-inner-dispatch in academic sizing tools is empirically robust on the question those tools answer — recommended capacity. Operational stochastic-dispatch (Tesla / Fluence / Wärtsilä) is independently justified by the realized-NPV uplift. The architectural divergence is empirically supported, and any operator can verify it for their market by running the diagnostic.

Pre-registered protocol locks all design choices before data touch (commits 2ce93ae, ebb035c, 0b5c9f1; see PREREGISTRATION_*.md).

2 Setup

2.1 Environment and dispatch

A single-battery hourly-arbitrage MDP. Continuous SoC; exogenous hourly price trace. Actions $a \in [-b_P, +b_P]$ (positive = discharge). SoC dynamics $\text{SoC}_{t+1} = \text{SoC}_t - a_t$ ($\eta = 1$). Per-step reward $r_t = p_t a_t - \alpha a_t^2$; the quadratic penalty is a convex surrogate for cycle depth, consistent with Shi et al. (2017).

Lifetime degradation post-hoc via rainflow counting (ASTM E1049) with stress function $f(\delta) = (k_1 \delta^{k_2} + k_3)^{-1}$ (Xu, 2018). Replacements scheduled when cumulative D consumes a 20% loss-of-health budget.

2.2 Price processes

Synthetic. Diurnal AR(1) hourly price (`price_signal.synth.diurnal`): mean 50, two diurnal harmonics (24 h, 12 h), AR(1) noise $\rho = 0.7, \sigma = 8$. Marginal mean 48.5, std 28.4.

Danish DK1. Energinet (Energinet, 2024) hourly DA spot prices, West Denmark, 2021/2022/2023.

Table 1: Market statistics, hourly day-ahead prices.

Market	Year	Currency	Mean	Std	Min	Max	Pers. residual std
DK1	2021	EUR/MWh	88.1	64.8	-43.9	620.0	46.4
DK1	2022	EUR/MWh	219.0	145.4	-19.0	871.0	84.8
DK1	2023	EUR/MWh	86.8	48.8	-440.1	524.3	43.3
ERCOT N.	2021	USD/MWh	146.6	892.6	0.0	8999.0	492.3
ERCOT N.	2022	USD/MWh	64.6	89.5	2.4	2543.8	92.8
ERCOT N.	2023	USD/MWh	55.5	218.7	1.1	4200.0	177.5

ERCOT North Hub. Annual day-ahead settlement-point prices via `gridstatus` (gridstatus.io, 2024)’s `get_dam_spp(year)` (free, public). 2021 contains February’s Storm Uri.

Statistics in Table 1. The two markets span very different regimes: DK1 is wind-heavy with multi-day weather-driven structure; ERCOT is thermal-dominant with rare-event price spikes. Both have years (DK1 2022, ERCOT 2021) with regime-defining shocks.

2.3 Forecasts and policies

Single forecast. Persistence at $t - 24$ h: $\hat{p}_t = p_{t-24}$. Standard baseline in price-forecasting literature; biased and fat-tailed.

Multi-lag ensemble ($K=4$). Persistence at lags $\{24, 48, 168, 336\}$ hours, averaged. Each member is a different hypothesis about which past pattern best predicts the present; averaging cancels lag-specific biases. Standard ensemble approach in load + price forecasting (Hong et al., 2014; Veron, 2014).

2×2 policy factorial. {linear LP via `scipy.linprog` HiGHS, quadratic QP via `cvxpy` CLARA-BEL} $\times$ {single, ensemble}, with cycling-cost coefficient $\alpha = 0.005$ (QP) or $\mu = 5.0$ /MWh (LP). Year-long horizons split into 8-week chunks; SoC carries forward across chunks.

2.4 Sizing model

Lifetime NPV = $R_{\text{annual}} \sum_{y=0}^{14} (1.07)^{-y} - C_E b_E - C_P b_P - \text{repl.}$ $C_E = 100\text{k currency-units} / \text{MWh}$, $C_P = 75\text{k currency-units} / \text{MW}$. 15-yr lifetime, 7% discount. Replacements pay only the energy component; PCS reused.

3 Theory

3.1 Heuristic proposition

Consider a price process $p(t)$ with spectral density $S_p(\omega)$. A pure tone at ω_0 , period $\tau_0 = 2\pi/\omega_0$, has dispatch-relevant capacity threshold $b_E^{\text{cyle}}(\tau_0) = b_P \tau_0 / 2$. Capacity past this contributes nothing to one-cycle revenue.

Define $q(\omega, \pi) \in [0, 1]$ as the fraction of price variance at frequency ω that policy π resolves from its forecast input. For a Gaussian-noise single forecast,

$$q(\omega, \pi_{\text{single}}) \approx \frac{S_p(\omega)}{S_p(\omega) + \sigma_f^2 / (1 - \rho_f^2)}$$

(Wiener-style SNR per frequency). Long timescales — low frequencies with signal power concentrated in fewer modes — are most sensitive to the noise floor. Define $\tau_{\text{res}}(\pi) = 2\pi/\omega_{\text{cut}}(\pi)$ where ω_{cut} is the lowest frequency at which q drops below threshold. We claim $b_{\text{sat}}(\pi) \approx b_P \tau_{\text{res}}(\pi)/2$. Argmax invariance follows when $\tau_{\text{res}}(\pi_1) = \tau_{\text{res}}(\pi_2)$.

Predicted regime characterization. Single-timescale processes (synthetic AR(1)): both policies resolve the dominant timescale; b_{sat} identical; argmax invariant. Multi-timescale processes (DK1 with weekly content; ERCOT with rare-event content): the deterministic policy may fail to resolve long-period arbitrage; b_{sat} differs across policies; argmax shifts.

The proposition is heuristic — the $q(\omega, \pi)$ form assumes Gaussian linearity, real forecast errors are fat-tailed and regime-conditional — but the qualitative prediction is the diagnostic test it scaffolds.

3.2 The b_{sat}^ϵ overlap diagnostic

For a given price process and dispatch policy, define the empirical revenue plateau as the smallest b_E at which marginal revenue $R'(b_E, \pi)$ falls below ϵ :

$$b_{\text{sat}}^\epsilon(\pi) = \min\{b_E : R'(b_E, \pi) < \epsilon\}.$$

Compute via finite differences on a fine b_E sweep. Bootstrap CIs across forecast realizations (or, with deterministic multi-lag ensembles, across overlapping 8-week subsamples).

The diagnostic test. Given two policies π_1, π_2 , the bootstrap CIs on $b_{\text{sat}}^\epsilon(\pi_1)$ and $b_{\text{sat}}^\epsilon(\pi_2)$ either overlap or disjoint. Pre-committed thresholds (PREREGISTRATION.ERCOT.md):

- *Invariance survives:* 95% CIs overlap by $\geq 50\%$ of the smaller CI's width.
- *Invariance breaks:* 95% CIs disjoint with a gap $\geq 25\%$ of the smaller CI's lower bound.
- *Inconclusive:* anything else.

Report at $\epsilon/C_E \in \{0.01, 0.05, 0.10\}$; headline at $\epsilon/C_E = 0.05$.

The diagnostic is independent of the proposition. The proposition predicts that single-timescale processes give overlap and multi-timescale processes give disjoint. The diagnostic just measures the overlap, on whatever data the practitioner has.

4 Results

4.1 Spectra and predicted divergence

Welch PSD on synthetic AR(1) returns dominant peaks at 23.5 h and 12.0 h with negligible weight at > 100 h (Figure 1). DK1 has strong peaks at 168 h (1 week, especially in 2022), 336 h (fortnight, 2023), 84 h, 224 h. ERCOT 2021 has broadband content from the Uri spike events. The proposition predicts invariance to hold on synthetic, break on DK1 and ERCOT.




Figure 1: Welch PSD: synthetic vs DK1 + ERCOT day-ahead. Both real markets have substantial multi-day spectral weight; the proposition predicts invariance breaking, but §4.3 falsifies the prediction.

Table 2: DK1 NPV at b_E^* , by year and policy. b_E^* identical across policies in every year.

Year	b_E^* (MWh)	NPV LP-single	NPV LP-ens	NPV QP-single	NPV QP-ens	
2021	4	247k	392k	310k	416k	(+34%)
2022	8	1.53M	1.86M	1.59M	1.91M	(+20%)
2023	4	386k	548k	462k	575k	(+25%)

4.2 Operational lift confirms the literature

Multi-lag ensemble dispatch yields positive realized-revenue lift over single-forecast dispatch in the relevant capacity range, on both markets. NPV uplift at the optimum spans +0.9% (ERCOT 2023, low-volatility regime) to +36.7% (ERCOT 2021 with Uri spikes), with DK1 in the middle (+20 to +34%). The wide range reflects how much exploitable forecast asymmetry each year contains. Consistent in sign and magnitude with [Cao et al. \(2020\)](#), [Krishnamurthy et al. \(2018\)](#).

4.3 Diagnostic returns “invariance survives” on every regime tested

4.4 Falsification of the proposition’s regime prediction

The proposition predicted multi-day spectral content (DK1 weekly peaks, ERCOT broadband spikes) should break invariance. It does not. The realized NPV curves are approximately vertical translations of each other in the relevant capacity range, not tilts.

This is a clean negative result for the regime-shift prediction and a positive result for the diagnostic: the test correctly returns “overlap”; the proposition that motivated it was wrong about which regimes would yield “disjoint”. The diagnostic survives because its job was to measure — not to predict — the overlap status of b_{sat}^ϵ CIs.

4.5 Charge/discharge pattern on the spike week

Figure 5 shows all four policies on the DK1 2022 spike week (containing the year’s max price 871/MWh on 2022-09-01, around a low-wind period visible in the resource panel). The dispatch decisions are visually nearly identical across policies, confirming the empirical similarity that the diagnostic measures.




Figure 2: Annual realized revenue vs b_E . Multi-lag ensemble dominates single-forecast in the relevant range. Past $b_E \sim 100$ MWh the curves cross — ensemble’s longer-period information is not net-positive at extreme capacity — but this region is far past NPV-optimal.

Table 3: ERCOT HB.NORTH NPV at b_E^* , by year and policy. b_E^* identical across policies in every year. NPV uplift at the optimum varies from +0.9% (2023) to +36.7% (2021 / Uri).

Year	b_E^* (MWh)	NPV LP-single	NPV LP-ens	NPV QP-single	NPV QP-ens	
2021 (Uri)	4	\$556k	\$776k	\$609k	\$832k	(+36.7%)
2022	8	\$788k	\$837k	\$840k	\$900k	(+7.1%)
2023	8	\$2.09M	\$2.13M	\$2.17M	\$2.19M	(+0.9%)

5 Discussion

5.1 Why the empirical answer is broader than the proposition predicted

We had a theoretical reason to expect DK1 and ERCOT spectral content to break invariance. They did not. Three candidate mechanisms we did not isolate:

- **Replacement scheduling absorbs the tilt.** Cycling depth differs by capacity in a way that changes the integer-valued replacement schedule across policies; the discrete steps push back against an otherwise-tilted lift.
- **Power-capacity bottleneck.** With $b_P = 1$ MW fixed, longer-period arbitrage requires longer to charge/discharge, capping the extra revenue larger batteries can extract. The multi-day arbitrage opportunities require slow accumulation that hits the b_P ceiling before the b_E ceiling.
- **Forecast resolution at long timescales is uniformly bad.** If both single and ensemble policies fail at multi-day timescales, neither captures multi-day arbitrage; bigger batteries help




Figure 3: Lifetime NPV vs b_E for DK1 by year. Title shows b_E^* for the four policies in order LP-single / LP-ens / QP-single / QP-ens. All four are identical in every panel. ERCOT panels analogous (Figure ??).

neither. Improving forecasts at long timescales (probabilistic ML weather forecasts) might expose a sizing shift this study did not detect.

5.2 Implication for sizing tools

Across the markets and regimes tested, deterministic-LP-inner-dispatch in academic sizing tools (hydesign, REopt, PyPSA) is empirically robust on capacity recommendations. Operational stochastic dispatch is independently justified by the 9–34% realized-NPV uplift at the recommended capacity. The architectural divergence between commercial operational platforms and academic sizing tools is empirically supported on both DK1 and ERCOT — not naive.

5.3 Limitations

1. One-dimensional sizing slice. b_E swept at fixed $b_P = 1$ MW; the proposition’s $b_{\text{sat}} \propto b_P$ form predicts that 2-D (b_E, b_P) surfaces could behave differently. Untested.
2. Multi-lag persistence is a weak ensemble. Modern probabilistic forecasts (deep quantile regression, ML weather) carry more information per member. We tested the cheapest defensible ensemble and acknowledge the negative result is conditional on this choice.
3. No scenario stochastic LP. Production-grade comparison (Birge & Louveaux, 1997; Krishnamurthy et al., 2018) would use 50–200 scenarios. We did not implement.
4. Replacement-count discreteness. Step changes can mask continuous shifts. Continuous-SoH model is future work.

Ensemble lift vs b_E . Approximately constant-in- $b_E \rightarrow$ argmax invariant.




Figure 4: Ensemble revenue lift over single, by year (DK1). Approximately uniform in b_E in the NPV-relevant range, supporting the empirical argmax-invariance. ERCOT panels analogous.

5. Battery-only. Co-located wind/PV with grid-cap-binding curtailment is a different problem; this paper isolates the price-driven arbitrage component.
6. Two markets (DK1, ERCOT). CAISO, PJM, and Nord Pool intraday have different price structure and could differ; the diagnostic generalises but the empirical conclusion is conditional on what we tested.

The most consequential limitation is (2): a richer ensemble might break invariance, in which case the diagnostic correctly fires under richer forecasts and the negative result is conditional on the cheap-ensemble baseline. We frame conservatively: *multi-lag persistence ensemble does not break invariance on DK1 or ERCOT, on the slice tested.*

6 Conclusion

We propose a regime-classification diagnostic for whether dispatch fidelity matters for battery sizing on a given market: the b_{sat}^e overlap test. The diagnostic is practitioner-runnable on one year of historical price data; it is falsifiable; and it survives independently of the theoretical proposition that motivates it. We test the diagnostic on synthetic, three years of DK1 (Energinet, including 2022 European energy crisis), and three years of ERCOT North Hub (gridstatus, including February 2021 Storm Uri). On every regime tested the diagnostic returns “invariance survives”: b_E^* is invariant across deterministic and stochastic dispatch, despite 9–34% realized-NPV uplift at the optimum. The proposition predicted regime-dependent invariance breaking on multi-day spectral content; the data falsify the prediction. The architectural divergence between commercial operational platforms (Tesla / Fluence / Wäertsilä, all stochastic) and academic sizing tools (hydesign / REopt / PyPSA, all deterministic LP inner) is empirically supported across the regimes tested.

Charge/discharge schedule at $b_E^* = 8.0$ MWh, $b_P = 1.0$ MW. Single-forecast policies react to phantom spikes; ensemble smooths.




Figure 5: DK1 2022 spike week. Top: realized DA price + multi-lag persistence forecasts + ensemble mean. Second: physical context (wind, solar, load, residual). Third: SoC trajectories under each of 4 policies (overlapping). Bottom: net dispatch (overlapping). Different policies make near-identical decisions on this week.

Reproducibility. Benchmark, four dispatch policies, multi-lag ensemble, sizing harness, pre-registration commits, and ERCOT + DK1 loaders are in `battery_gym` at https://github.com/kilojoules/battery_gym (commit hash in repo).

References

- Birge, J. R., Louveaux, F. (1997). *Introduction to Stochastic Programming*. Springer.
- Cao, J., Harrold, D., Fan, Z., Morstyn, T., Healey, D., Li, K. (2020). Deep Reinforcement Learning-Based Energy Storage Arbitrage With Accurate Lithium-Ion Battery Degradation Model. *IEEE Transactions on Smart Grid*, 11(5):4513–4521.
- Energinet (2024). Elspotprices public dataset. <https://www.energidaservice.dk/tso-electricity/elspotprices>.

- Fluence Energy (2024). Mosaic Intelligent Bidding Software. <https://fluenceenergy.com/mosaic-intelligent-bidding-software/>.
- gridstatus.io (2024). Open-source Python wrapper for ISO market data. <https://github.com/gridstatus/gridstatus>.
- Hong, T., Pinson, P., Fan, S. (2014). Global energy forecasting competition 2012. *International Journal of Forecasting*, 30(2):357–363.
- Krishnamurthy, D., Uckun, C., Zhou, Z., Thimmapuram, P., Botterud, A. (2018). Energy storage arbitrage under day-ahead and real-time price uncertainty. *IEEE Transactions on Power Systems*, 33(1):84–93.
- Murcia León, J. P., Habbou, H., Friis-Møller, M., Gupta, M., Zhu, R., Das, K. (2024). HyDesign: a tool for sizing optimization of grid-connected hybrid power plants. *Wind Energy Science*, 9:759–776.
- National Renewable Energy Laboratory (2025). Battery Degradation Modeling in Hybrid Power Plants. NREL/TP-FY25-84791.
- Shi, Y., Xu, B., Wang, D., Zhang, B. (2017). Using Battery Storage for Peak Shaving and Frequency Regulation: Joint Optimization for Superlinear Gains. *arXiv:1703.07968*.
- Tesla, Inc. (2024). Autobidder Support Documentation. <https://www.tesla.com/support/energy/tesla-software/autobidder>.
- Wärtsilä Energy (2025). GEMS Pulse press release.
- Weron, R. (2014). Electricity price forecasting: a review of the state-of-the-art. *International Journal of Forecasting*, 30(4):1030–1081.
- Xu, B. (2018). *Modeling and Control of Energy Storage Systems for Power System Operations*. PhD thesis, University of Washington.